

CSCI 455: Lab #4 — MPI Timing Model and Comparison

Darwin Jacob Groskleg

Winter 2020

Contents

Part 1: Estimating Communication Latency of MPI Send/Recv	1
mpi_latency.c	3
slr.h	7
slr.c	8
Part 2: The Efficient Processor Scaling of MPI_Bcast Over Send/Recv	10
Program Outputs and Results	10
compare_bcast.c	11
hiding_latency.c	13

Summary Lab 4 consists of a series of tests to measure the timing of a single send/recv communication using the ping-pong method. Based on the tests, you can estimate the t_{startup} and t_{data} with the least square regression method.

Part 1: Estimating Communication Latency of MPI Send/Recv

```

tags
darwin@groskleg@starbuck ~/Dropbox/Documents/Terms/2020-01 - Winter/CSC
I455/Lab4-Timing Model Comparison make run1
cc -std=c99 -Wall -Wextra -g -D_GLIBCXX_DEBUG -O0 -c -o slr.o slr.c
mpicc -std=c99 -Wall -Wextra -g -D_GLIBCXX_DEBUG -O0 -lmpi mpi_latency.c
slr.o -o mpi_latency
Platform: Darwin (4 cpu cores recognized)
PMIX_MCA_gds=hash mpirun --host localhost --mca btl_vader_backing_directory
/tmp --mca btl ^tcp --oversubscribe -np 4 ./mpi_latency
INFO: Number of processes = 4
INFO: Only executing 2 tasks - extra cluster processes will be ignored
task 1 has started...
task 0 has started...
Beginning latency timing test:
    Number of reps = 10
    Data Size      = 10
*****
Rep#      T0      T1      deltaT
  1         0.00    52.00    52.00
  2        60.00    86.00    26.00
  3        92.00   104.00    12.00
  4       109.00   112.00     3.00
  5       116.00   119.00     3.00
  6       123.00   125.00     2.00
  7       129.00   132.00     3.00
  8       135.00   137.00     2.00
  9       141.00   143.00     2.00
 10       148.00   150.00     2.00
*****
*** Avg round trip time = 10.700000 microseconds
*** Avg one way latency = 5.350000 microseconds
Beginning latency timing test:
    Number of reps = 10
    Data Size      = 50
*****

```

Figure 1: First part of output

```
Beginning latency timing test:
  Number of reps = 10
  Data Size      = 5000
*****
Rep#      T0      T1      deltaT
  1      1104.00    1451.00    347.00
  2      1458.00    1473.00     15.00
  3      1483.00    1494.00     11.00
  4      1498.00    1505.00      7.00
  5      1509.00    1517.00      8.00
  6      1521.00    1529.00      8.00
  7      1533.00    1540.00      7.00
  8      1544.00    1552.00      8.00
  9      1555.00    1563.00      8.00
 10      1567.00    1575.00      8.00
*****
*** Avg round trip time = 42.700000 microseconds
*** Avg one way latency = 21.350000 microseconds
SUMMARY STATISTICS/ESTIMATIONS:
t_comm(1) = 5.060000 microseconds
t_startup = 5.060000 microseconds
t_data    = -0.000000 microseconds
darwingroskleg@starbuck > ~/Dropbox/Documents/Terms/2020-01 - Winter/CSC
darwingroskleg@starbuck > ~/Dropbox/Documents/Terms/2020-01 - Winter/CSC
I455/Lab4-Timing Model Comparison >
```

Figure 2: Last part of output

mpi_latency.c

```

1  /*****
2  * FILE: mpi_latency.c
3  * AUTHORS:
4  *   Darwin Jacob Groskleg (2020)
5  *   Sazzad (02/11/18)
6  *   Laurence T. Yang
7  * DESCRIPTION:
8  *   MPI Latency Timing Program – C Version
9  *   In this example code, a MPI communication timing test is performed.
10 *   MPI process 0 will send "reps" number of 1 byte messages to MPI process 1,
11 *   waiting for a reply between each rep. Before and after timings are made
12 *   for each rep and an average calculated when completed.
13 *
14 * NOTES
15 *   - Ping-pong method sends same size data back and forth.
16 *****/
17
18 #include <stdio.h>
19 #include <stdlib.h>
20 #include <time.h>
21 #include <string.h>
22 #include <stdbool.h>
23 #include <math.h>
24 #include <assert.h>
25
26 #include <sys/time.h>
27 #include <mpi.h>
28
29 #include "slr.h"
30
31 #define NUMBER_REPS 10
32 #define DATA_SIZE 5000
33
34 /* send/receive process designators, maps task label to expected rank */
35 enum TaskRanks {
36     Master = 0,
37     Worker = 1
38 };
39
40 double sample_latency_with_n_bytes(int bytes_of_traffic);
41 void run_sampling_responder(int bytes_of_traffic);
42
43 /*
44 * Estimate the t_startup and t_data with the least square regression method:
45 * y: AvgT/2 (one way latency)
46 * let K = 10 (number of runs, different message sizes)
47 * t_startup = m * 0 + b (time to send msg with no data)
48 * t_comm = m*1 + b (time to startup and send 1 data word)
49 * t_data = t_comm - t_startup (time to send one data word)
50 */
51 int main (int argc, char *argv[]) {
52     MPI_Init(&argc, &argv);
53     int cluster_size; /* number of MPI processes */
54     MPI_Comm_size(MPI_COMM_WORLD, &cluster_size);
55     int rank; /* my MPI process number */

```

```

56 MPI_Comm_rank(MPI_COMM_WORLD, &rank);
57
58 if (rank == Master && cluster_size != 2) {
59     printf("INFO: Number of processes = %d\n", cluster_size);
60     printf("INFO: Only executing 2 tasks - extra cluster processes ");
61     printf("will be ignored\n");
62 }
63 MPI_Barrier(MPI_COMM_WORLD);
64
65 if (rank < 2) {
66     printf("task %d has started...\n", rank);
67 }
68 MPI_Barrier(MPI_COMM_WORLD);
69
70 int x_byte_counts[] = {10, 50, 100, 200, 500, 1000, 2000, 3000, 4000, 5000};
71 size_t k_tests = sizeof(x_byte_counts)/sizeof(x_byte_counts[0]);
72 double y_timings[k_tests];
73
74 for (size_t i=0; i<k_tests; i++){
75     if (rank == Worker) {
76         run_sampling_responder(x_byte_counts[i]);
77     }
78     else if (rank == Master) {
79         y_timings[i] = sample_latency_with_n_bytes(x_byte_counts[i]);
80     }
81 }
82
83 if (rank == Master) {
84     slr_equation_t eqt = slr_find_line(k_tests, x_byte_counts, y_timings);
85     double t_startup = slr_predict(eqt, 0);
86     double t_comm = slr_predict(eqt, 1);
87     double t_data = t_comm - t_startup;
88     printf("SUMMARY STATISTICS/ESTIMATIONS:\n");
89     printf("t_comm(1) = %f microseconds\n", t_comm);
90     printf("t_startup = %f microseconds\n", t_startup);
91     printf("t_data = %f microseconds\n", t_data);
92 }
93
94 MPI_Finalize();
95 exit(0);
96 }
97
98 /* Returns: time in microseconds (t_comm)
99 * - the one-way latency (round-trip / 2)
100 * Takes:
101 * bytes_of_traffic = N, the amount of data to be sent
102 */
103 double sample_latency_with_n_bytes(int bytes_of_traffic) {
104     assert(bytes_of_traffic <= DATA_SIZE);
105
106     char msg[DATA_SIZE]; /* buffer containing DATA_SIZE byte message */
107     MPI_Status status; /* MPI receive routine parameter */
108     int tag = 1; /* MPI message tag parameter */
109     int reps = NUMBER_REPS; /* number of samples per test */
110     /* round-trip latency timing test */
111     printf("Beginning latency timing test:\n");

```

```

112     printf("\tNumber of reps = %d\n", reps);
113     printf("\tData Size      = %d\n", bytes_of_traffic);
114     printf("*****\n");
115     /*      < > < . > < . > < . >      */
116     printf("Rep#      T0      T1      deltaT\n");
117
118     const double MSecond = pow(10, 6);
119     const double ClockResolution = 1; //MPI_Wtick(); // seconds per clock tick
120     double T0, T1;                /* start/end times per rep in ms */
121     double deltaT;                /* time for one rep in ms */
122     double sumT = 0;              /* sum of all reps times in microseconds */
123     int error;
124     for (int n = 1; n <= reps; n++) {
125         /* start time */
126         T0 = MPI_Wtime() * MSecond * ClockResolution;
127
128         /* send message to worker - message tag set to 1. */
129         error = MPI_Send(
130             &msg,
131             bytes_of_traffic,
132             MPI_BYTE,    // for char
133             Worker,      // destination
134             tag,
135             MPI_COMM_WORLD
136         );
137         /* If return code indicates error quit */
138         // DARWIN GROSLEG:
139         // this is unnecessary and already properly handled
140         // by the MPI error handler, which will abort the MPI job in most
141         // cases.
142         if (error)
143             MPI_Abort(MPI_COMM_WORLD, error);
144
145         /* Now wait to receive the echo reply from the worker */
146         // "echoes" the same data back???
147         error = MPI_Recv(
148             &msg,
149             bytes_of_traffic,
150             MPI_BYTE,    // for char
151             Worker,      // source
152             tag,
153             MPI_COMM_WORLD,
154             &status);
155
156         /* If return code indicates error quit */
157         // Redundant step is skipped because the MPI already handles it.
158         if (error)
159             MPI_Abort(MPI_COMM_WORLD, error);
160
161         /* end time */
162         T1 = MPI_Wtime() * MSecond * ClockResolution;
163
164         /* calculate round trip time and print */
165         deltaT = T1 - T0;
166         sumT += deltaT;
167         /* print statement for each to keep each column right justified */

```

```

168     printf("%4d ", n);
169     printf("%10.2f ", T0);
170     printf("%10.2f ", T1);
171     printf("%10.2f\n", deltaT);
172 }
173
174 /* average time per rep in microseconds */
175 double avgT = sumT / reps;
176 printf("*****\n");
177 printf("*** Avg round trip time = %f microseconds\n", avgT);
178 printf("*** Avg one way latency = %f microseconds\n", avgT/2);
179
180 return avgT/2;
181 }
182
183 void run_sampling_responder(int bytes_of_traffic) {
184     assert(bytes_of_traffic <= DATA_SIZE);
185     char msg[DATA_SIZE]; /* buffer containing DATA_SIZE byte message */
186     MPI_Status status; /* MPI receive routine parameter */
187     int tag = 1; /* MPI message tag parameter */
188     int reps = NUMBER_REPS; /* number of samples per test */
189
190     while (reps--) {
191         // ping
192         MPI_Recv(
193             &msg,
194             bytes_of_traffic,
195             MPI_BYTE,
196             Master,
197             tag,
198             MPI_COMM_WORLD,
199             &status);
200         // pong
201         MPI_Send(
202             &msg,
203             bytes_of_traffic,
204             MPI_BYTE,
205             Master,
206             tag,
207             MPI_COMM_WORLD);
208     }
209 }

```

slr.h

```
1  /* slr.h
2  * -----
3  * Authors: Darwin Jacob Groskleg (2020)
4  *
5  * Purpose: do simple linear regression.
6  */
7  #ifndef SLR_H_INCLUDED
8  #define SLR_H_INCLUDED
9
10 typedef struct {
11     double a; /* intercept */
12     double b; /* slope */
13 } slr_equation_t;
14
15 slr_equation_t slr_find_line(int n, int X[], double Y[]);
16
17 double slr_predict(slr_equation_t eqt, int x);
18
19 #endif /* SLR_H_INCLUDED */
```


slr.c

```

1  /* slr.c
2  * -----
3  * Authors: Darwin Jacob Groskleg (2020)
4  *
5  * Purpose: do simple linear regression.
6  */
7  #include "slr.h"
8
9  #include <assert.h>
10 #include <math.h>
11
12 double regression_coefficient(int, double, double, double, double);
13 double y_intercept(int, double, double, double);
14
15
16 double slr_predict(slr_equation_t eqt, int x) {
17     /* intercept + slope * x */
18     return eqt.a + eqt.b * x;
19 }
20
21 /* let K = n
22 * for each (x, y) of K:
23 *     x^2, xy
24 * calc slope:
25 *     m = (K * SUM(xy) - SUM(x)*SUM(y))
26 *         / (K * SUM(x^2) - SUM(x)^2)
27 * calc intercept:
28 *     b = (SUM(y) - m*SUM(x))/K
29 *
30 * passes as pointer, NOT COPY
31 */
32 slr_equation_t slr_find_line(int n, int X[], double Y[]) {
33     //assert() len(X) == len(Y) == n
34     slr_equation_t eqt;
35     double sigma_x = 0;
36     double sigma_xx = 0;
37     double sigma_y = 0;
38     double sigma_xy = 0;
39     for (int i=0; i<n; i++) {
40         sigma_x += X[i];
41         sigma_xx += pow(X[i], 2);
42         sigma_y += Y[i];
43         sigma_xy += X[i] * Y[i];
44     }
45     double slope = regression_coefficient(n, sigma_x, sigma_xx,
46                                         sigma_y, sigma_xy);
47     eqt.a = y_intercept(n, slope, sigma_x, sigma_y);
48     eqt.b = slope;
49     return eqt;
50 }
51
52 /* Slope aka "Regression Coefficient"
53 *
54 * Biostatistical Analysis 5ed, JH Zar, Pages 330-337
55 *  $Y_i = a + BX_i$  for best fit using least squares linear regression.

```

```
56  *
57  */
58  double regression_coefficient(
59      int    data_points,
60      double sum_of_x,      // SUM x
61      double sum_sqr_of_x,  // SUM xx
62      double sum_of_y,      // SUM y
63      double sum_of_xy)     // SUM xy
64  {
65      double sum_of_cross_products = sum_of_xy - sum_of_x * sum_of_y/data_points;
66      double sum_of_squares_x = sum_sqr_of_x - pow(sum_sqr_of_x, 2)/data_points;
67      return sum_of_cross_products / sum_of_squares_x;
68  }
69
70  double y_intercept(
71      int    data_points,
72      double slope,
73      double sum_of_x,
74      double sum_of_y)
75  {
76      return (sum_of_y - slope * sum_of_x) / data_points;
77  }
```

Part 2: The Efficient Processor Scaling of MPI_Bcast Over Send/Recv

Program Outputs and Results

compare_bcast.c

```

1  //
2  // Comparison of MPI_Bcast with the my_bcast function
3  //
4  #include <stdio.h>
5  #include <stdlib.h>
6  #include <assert.h>
7
8  #include <mpi.h>
9
10 /* send/receive process designators, maps task label to expected rank */
11 enum TaskRanks {
12     Master = 0,
13     Worker = 1
14 };
15
16 void my_bcast(void* data, int count, MPI_Datatype datatype, int root,
17              MPI_Comm communicator)
18 {
19     int world_rank;
20     MPI_Comm_rank(communicator, &world_rank);
21     int world_size;
22     MPI_Comm_size(communicator, &world_size);
23
24     if (world_rank == root) {
25         // If we are the root process, send our data to everyone
26
27     } else {
28         // If we are a receiver process, receive the data from the root
29
30     }
31
32 }
33
34
35 int main(int argc, char** argv) {
36     if (argc != 3) {
37         fprintf(stderr, "Usage: compare_bcast num_elements num_trials\n");
38         exit(1);
39     }
40
41     int num_elements = atoi(argv[1]);
42     int num_trials = atoi(argv[2]);
43
44     MPI_Init(NULL, NULL);
45
46     int world_rank;
47     MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);
48
49     double total_my_bcast_time = 0.0;
50     double total_mpi_bcast_time = 0.0;
51     int i;
52     int* data = (int*)malloc(sizeof(int) * num_elements);
53     assert(data != NULL);
54
55     for (i = 0; i < num_trials; i++) {

```

```
56     // Time my_bcast
57     // Synchronize before starting timing
58
59
60
61     // Synchronize again before obtaining final time
62
63
64
65     // Time MPI_Bcast
66
67
68
69 }
70
71 // Print off timing information
72 if (world_rank == 0) {
73     printf("Data size = %d, Trials = %d\n", num_elements * (int)sizeof(int),
74           num_trials);
75     printf("Avg my_bcast time = %lf\n", total_my_bcast_time / num_trials);
76     printf("Avg MPI_Bcast time = %lf\n", total_mpi_bcast_time / num_trials);
77 }
78
79 free(data);
80 MPI_Finalize();
81 return 0;
82 }
```

hiding_latency.c

```

1  /* hiding_latency.c
2  * -----
3  * Authors: Darwin Jacob Groskleg
4  * Date:    Saturday, May 16, 2020
5  *
6  * Taken from "Intro to Parallel Computing (2018), Roman Trobec et al.",
7  * Page 120
8  *
9  * Overlapping communication and computation
10 */
11 #include <mpi.h>
12 #include <stdlib.h>
13 #include <math.h>
14 #include <stdio.h>
15 int i; double a;
16 for (i = 0; i < 100000000/numproc; i++) {
17     a = sin(sqrt(i)); //different amount of calculation return a;
18     main(int argc, char* argv[]) //number of processes must be > 1
19     int p, i, myid, tag=1, proc, ierr;
20     double start_p, run_time, start_c, comm_t, start_w, work_t, work_r; double *buff = nullptr;
21     MPI_Request request; MPI_Status status;
22     MPI_Init(&argc, &argv);
23     start_p = MPI_Wtime(); MPI_Comm_rank(MPI_COMM_WORLD, &myid); MPI_Comm_size(MPI_COMM_WORLD, &p);
24     #define master 0
25     #define MSGSIZE 100000000 //5000000 //different sizes of ←
26     messages
27     buff = (double*)malloc(MSGSIZE * sizeof(double)); //allocate
28     if (myid == master) {
29         for (i = 0; i < MSGSIZE; i++) { //initialize message
30             buff[i] = 1;
31         }
32         start_c = MPI_Wtime();
33         for (proc = 1; proc<p; proc++) {
34             MPI_Irecv(buff, MSGSIZE,
35             MPI_DOUBLE, MPI_ANY_SOURCE, tag, MPI_COMM_WORLD, &
36                 1
37             2
38             3
39             4 5{ 6
40                 double other_work(int numproc)
41                 7
42             8 9}
43                 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28
44                 29 30 31 32 33 34 35 36 37 38 39
45                 40 41 42 43
46                 44 45 46 47 48 49 50 51
47             }
48             int {
49                                     #if 1
50                 //non-blocking receive
51                 request); #endif
52                 #if 0
53             ); #endif
54         }

```

```
55 | MPI_Recv(buff, MSGSIZE, //blocking receive
56 | MPI_DOUBLE, MPI_ANY_SOURCE, tag, MPI_COMM_WORLD, &status←
57 |     }
58 | comm_t = MPI_Wtime() - start_c; start_w = MPI_Wtime();
59 | work_r = other_work(p);
60 | work_t = MPI_Wtime() - start_w; MPI_Wait(&request, &status);
61 | //block until Irecv is done
62 |
```