

# CSCI 455: Project 1 (Lab #5–6) — Image Transformation

Darwin Jacob Groskleg

Winter 2020

## Contents

<b>Program Output</b>	<b>1</b>
Console . . . . .	1
Filtered Images . . . . .	3
<b>Questions</b>	<b>4</b>
Q1. How does the single threaded perform in the serial vs parallel code? . . . . .	4
Q2. Which method of communication was implemented and why? . . . . .	4
Q3. How does the parallel code scale with number of processors? . . . . .	4
<b>Further Improvements</b>	<b>5</b>
<b>Code</b>	<b>5</b>
src/transform-parall.c . . . . .	6
src/filter_ppm.h . . . . .	10
src/filter_ppm.c . . . . .	11
src/ppm.h . . . . .	16
src/ppm.c . . . . .	17
src/transform-serial.c . . . . .	20
Makefile . . . . .	22

## Program Output

### Console

```
transcript written on tmp/Lab5-6-HANDIN.log.
darwingroskleg@starbuck ~/Dropbox/Documents/Terms/2020-01 - Winter/CSCI455/Lab5-6-Project_1_Image_Transformation ➤ make test-serial
mpicc -std=c99 -Wall -Wextra -g -D_GLIBCXX_DEBUG -O0 -c -o transform-serial.o ./src/transform-serial.c
mpicc transform-serial.o ppm.o filter_ppm.o -lm -o transform-serial
./transform-serial
infile name: ./im1.ppm
Neighborhood pixel radius |x y|: 5 5
Average intensity is: 108.889805
darwingroskleg@starbuck ~/Dropbox/Documents/Terms/2020-01 - Winter/CSCI455/Lab5-6-Project_1_Image_Transformation ➤
```

Figure 1: Output from serial implementation found in `src/transform-serial.c` for the first image.

```
darwingroskleg@starbuck ~/Dropbox/Documents/Terms/2020-01 - Winter/CSCI455/Lab5-6-Project_1_Image_Transformation ➤ make
mpicc -std=c99 -Wall -Wextra -g -D_GLIBCXX_DEBUG -O0 -c -o transform-serial.o ./src/transform-serial.c
mpicc -std=c99 -Wall -Wextra -g -D_GLIBCXX_DEBUG -O0 -c -o ppm.o ./src/ppm.c
mpicc -std=c99 -Wall -Wextra -g -D_GLIBCXX_DEBUG -O0 -c -o filter_ppm.o ./src/filter_ppm.c
mpicc -std=c99 -Wall -Wextra -g -D_GLIBCXX_DEBUG -O0 -c -o transform-parall.o ./src/transform-parall.c
mpicc -std=c99 -Wall -Wextra -g -D_GLIBCXX_DEBUG -O0 -c -o test_filter.o ./src/test_filter.c
mpicc transform-serial.o ppm.o filter_ppm.o -lm -o transform-serial
mpicc transform-parall.o ppm.o filter_ppm.o -lm -o transform-parall
mpicc test_filter.o ppm.o filter_ppm.o -lm -o test_filter
darwingroskleg@starbuck ~/Dropbox/Documents/Terms/2020-01 - Winter/CSCI455/Lab5-6-Project_1_Image_Transformation ➤ make test-parall-all
Platform: Darwin (4 cpu cores recognized)
PMIX_MCA_gds=hash mpirun --host localhost --mca btl_vader_backing_directory /tmp --mca btl ^tcp --oversubscribe -np 4 ./transform-parall im1.ppm
[0]: infile name: im1.ppm
[0]: file read, 515788 pixels, 676px wide, 763px in height
[0]: Averaging Filter 4 node parallel computation time: 8.8748e-08 seconds
[0]: Threshold Filter 4 node parallel computation time: 5.341e-09 seconds
[0]: Overall average intensity is: 108.889805
Platform: Darwin (4 cpu cores recognized)
PMIX_MCA_gds=hash mpirun --host localhost --mca btl_vader_backing_directory /tmp --mca btl ^tcp --oversubscribe -np 4 ./transform-parall im2.ppm
[0]: infile name: im2.ppm
[0]: file read, 486168 pixels, 590px wide, 824px in height
[0]: Averaging Filter 4 node parallel computation time: 8.8676e-08 seconds
[0]: Threshold Filter 4 node parallel computation time: 3.389e-09 seconds
[0]: Overall average intensity is: 43.199266
Platform: Darwin (4 cpu cores recognized)
PMIX_MCA_gds=hash mpirun --host localhost --mca btl_vader_backing_directory /tmp --mca btl ^tcp --oversubscribe -np 4 ./transform-parall im3.ppm
[0]: infile name: im3.ppm
[0]: file read, 419175 pixels, 575px wide, 729px in height
[0]: Averaging Filter 4 node parallel computation time: 7.4068e-08 seconds
[0]: Threshold Filter 4 node parallel computation time: 4.194e-09 seconds
[0]: Overall average intensity is: 50.299069
```

Figure 2: Console output from running parallel implementation found in `src/transform-parall.c` for all 3 image files.

```

mpicc -std=c99 -Wall -Wextra -g -O GLIBCXX_DEBUG -O0 -c -o transform-parall.o ./src/transform-parall.c
mpicc -std=c99 -Wall -Wextra -g -O GLIBCXX_DEBUG -O0 -c -o ppm.o ./src/ppm.c
mpicc -transform-parall.o ppm.o filter_ppm.o -lm -o transform-parall
time PMIX_MCA_gds=hash mpirun --host localhost --mca btl_vader_backing_directory /tmp --mca btl ^tcp --oversubscribe -np 1 ./transform-parall im1.ppm
[0]: infile name: im1.ppm
[0]: file read, 515788 pixels, 676px wide, 763px in height
[0]: Averaging Filter: Neighborhood radius |x y|: 5 5
[0]: Averaging Filter: 1 node parallel computation time: 3.27512e-07 seconds
[0]: Threshold Filter: Overall average intensity is: 108.889805
[0]: Threshold Filter: 1 node parallel computation time: 1.222e-08 seconds

real    0m0.465s
user    0m0.431s
sys     0m0.031s
time PMIX_MCA_gds=hash mpirun --host localhost --mca btl_vader_backing_directory /tmp --mca btl ^tcp --oversubscribe -np 2 ./transform-parall im1.ppm
[0]: infile name: im1.ppm
[0]: file read, 515788 pixels, 676px wide, 763px in height
[0]: Averaging Filter: Neighborhood radius |x y|: 5 5
[0]: Averaging Filter: 2 node parallel computation time: 1.67141e-07 seconds
[0]: Threshold Filter: Overall average intensity is: 108.889805
[0]: Threshold Filter: 2 node parallel computation time: 6.652e-09 seconds

real    0m0.512s
user    0m0.498s
sys     0m0.052s
time PMIX_MCA_gds=hash mpirun --host localhost --mca btl_vader_backing_directory /tmp --mca btl ^tcp --oversubscribe -np 4 ./transform-parall im1.ppm
[0]: infile name: im1.ppm
[0]: file read, 515788 pixels, 676px wide, 763px in height
[0]: Averaging Filter: Neighborhood radius |x y|: 5 5
[0]: Averaging Filter: 4 node parallel computation time: 1.07571e-06 seconds
[0]: Threshold Filter: Overall average intensity is: 108.889805
[0]: Threshold Filter: 4 node parallel computation time: 3.825e-09 seconds

real    0m0.233s
user    0m0.620s
sys     0m0.070s
dawringroskleg@Starbuck: ~/Dropbox/Documents/Terms/2020-01 - Winter/CSCI455/Lab5-6-Project_1_Image_Transformation>

```

Figure 3: Console output from node scaling performance comparison (see Q3).

```

sys    0m0.169s
~/Dropbox/Documents/Terms/2020-01 - Winter/CSCI455/Lab5-6-Project_1_Image_Transformation> make ql-perf
time ./transform-serial
infile name: ./im1.ppm
Neighborhood pixel radius |x y|: 5 5
Average intensity is: 108.889805
    1.00 real      0.91 user      0.06 sys
time PMIX_MCA_gds=hash mpirun --host localhost --mca btl_vader_backing_directory /tmp --mca btl ^tcp --oversubscribe -np 1 ./transform-parall
[0]: infile name: ./im1.ppm
[0]: file read, 515788 pixels, 676px wide, 763px in height
[0]: Averaging Filter: Neighborhood radius |x y|: 5 5
[0]: Averaging Filter: 1 node parallel computation time: 1.07571e-06 seconds
[0]: Threshold Filter: Overall average intensity is: 108.889805
[0]: Threshold Filter: 1 node parallel computation time: 6.0398e-08 seconds

real    0m1.477s
user    0m1.056s
sys     0m0.138s
~/Dropbox/Documents/Terms/2020-01 - Winter/CSCI455/Lab5-6-Project_1_Image_Transformation> make ql-perf
time ./transform-serial
infile name: ./im1.ppm
Neighborhood pixel radius |x y|: 5 5
Average intensity is: 108.889805
    1.23 real      0.91 user      0.09 sys
time PMIX_MCA_gds=hash mpirun --host localhost --mca btl_vader_backing_directory /tmp --mca btl ^tcp --oversubscribe -np 1 ./transform-parall
[0]: infile name: ./im1.ppm
[0]: file read, 515788 pixels, 676px wide, 763px in height
[0]: Averaging Filter: Neighborhood radius |x y|: 5 5
[0]: Averaging Filter: 1 node parallel computation time: 1.16512e-06 seconds
[0]: Threshold Filter: Overall average intensity is: 108.889805
[0]: Threshold Filter: 1 node parallel computation time: 5.2263e-08 seconds

real    0m1.553s
user    0m1.056s
sys     0m0.153s
~/Dropbox/Documents/Terms/2020-01 - Winter/CSCI455/Lab5-6-Project_1_Image_Transformation> make ql-perf
time ./transform-serial
infile name: ./im1.ppm
Neighborhood pixel radius |x y|: 5 5
Average intensity is: 108.889805
    1.15 real      0.91 user      0.06 sys
time PMIX_MCA_gds=hash mpirun --host localhost --mca btl_vader_backing_directory /tmp --mca btl ^tcp --oversubscribe -np 1 ./transform-parall
[0]: infile name: ./im1.ppm
[0]: file read, 515788 pixels, 676px wide, 763px in height
[0]: Averaging Filter: Neighborhood radius |x y|: 5 5
[0]: Averaging Filter: 1 node parallel computation time: 1.10525e-06 seconds
[0]: Threshold Filter: Overall average intensity is: 108.889805
[0]: Threshold Filter: 1 node parallel computation time: 6.9013e-08 seconds

real    0m1.500s
user    0m1.051s
sys     0m0.150s

```

Figure 4: Console output from single threaded performance comparison (see Q1).

## Filtered Images

Images shown before and after filtering using parallel implementation. Ordered original, fuzzy then threshold from the left. Averaging filter used a radius of (x=5, y=5).



Figure 5: From im1.ppm of *Small Cowper Madonna* by Raphael, 1505. Overall intensity of 108.889805.

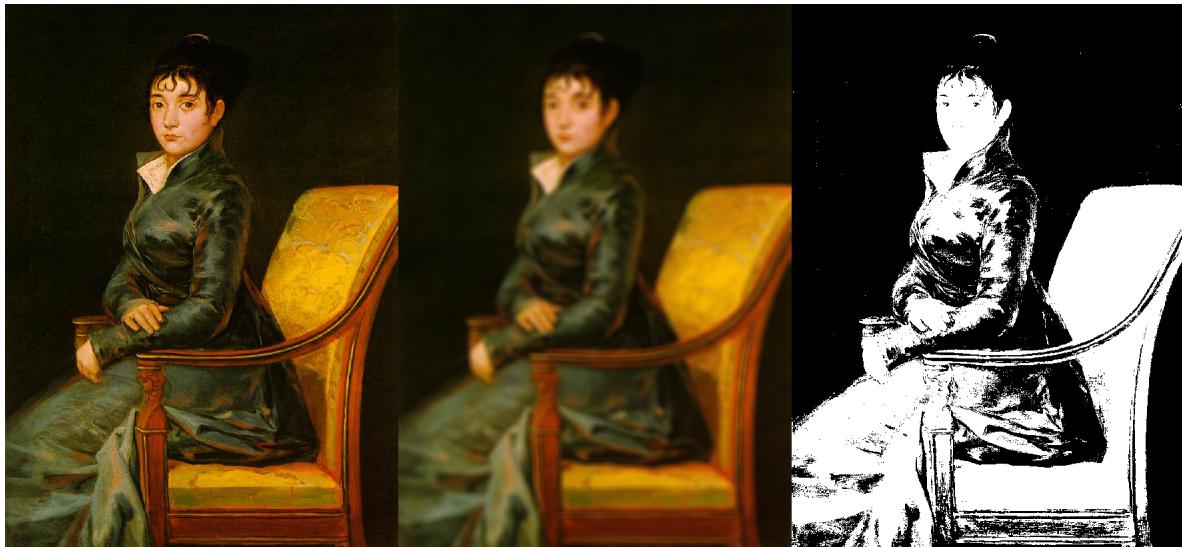


Figure 6: From im2.ppm of *Doña Teresa Sureda* by Francisco Goya, 1805. Overall intensity of 43.199266.



Figure 7: From im3.ppm of *Sir Brian Tuke* by Hans Holbein the Younger, c. 1532-1534. Overall intensity of 50.299069.

## Questions

### **Q1. How does the single threaded perform in the serial vs parallel code?**

Using the `time` command running both parallel (1 node) and serial on `im1.ppm` we get the following averages for `real` execution time in seconds:

- Serial: 1.15 seconds (1.06, 1.23, 1.15)
- Parallel: 1.51 seconds (1.477, 1.553, 1.500)

Four runs of the code in a row were made, throwing away the first where the data may not have been in the disk cache (mitigates outliers). The resulting performance shows the serial code, without the extra unused machinery of MPI to setup, is faster than the single node execution of the parallel code.

### **Q2. Which method of communication was implemented and why?**

There are 3 areas of communication: loading the initial image file data, running the fuzzy transformation and running the threshold transformation. First, opening the image file is done by the root process which then broadcasts the image data to all other processes. The fuzzy transformation only needs to communicate once, that is a Gather from all processes, their respective subset of the image data, to the root node. Finally, the threshold transformation communicates using All-Reduce in the calculation of the overall image intensity and then again with a Gather in the same way as the fuzzy transformation did.

### **Q3. How does the parallel code scale with number of processors?**

Program performance indexed by node count, tracked using `time`:

- 1: real 0.465s

- 2: real 0.512s
- 4: real 0.233s

Below are the performance times of the individual transform algorithms as measured in the program using MPI's given performance functions `WTick()` and `WTime()`.

The “fuzzy transformation” scaled as follows:

- 1: 3.27512 e-07 seconds
- 2: 1.67141 e-07 seconds
- 4: 0.87628 e-07 seconds

The “threshold transformation” scaled as follows:

- 1: 1.222 e-08 seconds
- 2: 0.6652 e-08 seconds
- 4: 0.3825 e-08 seconds

## Further Improvements

Improvements that were not required but I believe would yield better performance or speedup.

- Parallel I/O: use the MPI file I/O calls to read and write to the PPM image files. This would allow the expensive overhead of broadcasting (communicating) all the file data from the root process to the others. Especially true for the Averaging Filter. Most importantly this improves overall program performance.
- Pipe-lining with Cartesian Grid for communicating. Avoid calculating sums twice with the averaging filter.

## Code

See `src/transform-serial.c` for serial implementation (modified from prof. L.T. Yang). See `src/transform-parall.c` for the MPI/parallel implementation. All other sources files are used in both implementations. The `Makefile` is included to clarify how performance tests may have been conducted but personal/private makefiles libraries are not included in this document.

## src/transform-parall.c

```

1  /* transform-parall.c
2  * -----
3  * Authors: Darwin Jacob Groskleg
4  */
5  #include <assert.h>
6  #include <stdlib.h>
7  #include <string.h>
8  #include <stdio.h>
9
10 #include <mpi.h>
11
12 #include "ppm.h"
13 #include "filter_ppm.h"
14
15 enum TaskRanks {
16     ConsoleIORank = 0
17 };
18
19 void make_ppm_pixel_dt(MPI_Datatype* MPI_PIXEL);
20 char* append_to_basename(const char* filename, const char* appendix);
21
22 int main(int argc, char *argv[]) {
23     MPI_Init(&argc, &argv);
24     int rank;
25     MPI_Comm_rank(MPI_COMM_WORLD, &rank);
26     int cluster_size;
27     MPI_Comm_size(MPI_COMM_WORLD, &cluster_size);
28
29     const char* default_image_filename = "./im1.ppm";
30     char* image_filename;
31     if (argc > 1) {
32         image_filename = (char *) malloc( sizeof (char) * strlen(argv[1]));
33         strcpy(image_filename, argv[1]);
34     } else {
35         size_t default_len = strlen(default_image_filename);
36         image_filename = (char *) malloc( sizeof (char) * default_len);
37         strcpy(image_filename, default_image_filename);
38     }
39     if (rank == ConsoleIORank)
40         printf("[%d]: infile name: %s\n", rank, image_filename);
41
42     /* Create Pixel datatype for MPI */
43     MPI_Datatype MPI_PIXEL;
44     make_ppm_pixel_dt(&MPI_PIXEL);
45
46     double t0, tf, t_delta_seconds;
47     double seconds_per_tick = MPI_Wtick();
48
49 /*
50 * Read and Share Initial PPM Image Data
51 *
52 */
53     Pixel* in_image = (Pixel *) malloc(PPM_MAX_PIXELS * sizeof (Pixel));
54     assert(in_image != 0);
55

```

```

56     int width, height, max;
57     int err = 0;
58     if (rank == ConsoleIORank) {
59         err = read_ppm(image_filename, &width, &height, &max, (char *)in_image);
60         if (err != 0) MPI_Abort(MPI_COMM_WORLD, err);
61         printf("[%d]: file read, %d pixels, %dpx wide, %dpx in height\n",
62                rank, width*height, width, height);
63     }
64     MPI_Bcast(&width, 1, MPI_INT, ConsoleIORank, MPI_COMM_WORLD);
65     MPI_Bcast(&height, 1, MPI_INT, ConsoleIORank, MPI_COMM_WORLD);
66     MPI_Bcast(&max, 1, MPI_INT, ConsoleIORank, MPI_COMM_WORLD);
67     MPI_Bcast(in_image, width*height, MPI_PIXEL, ConsoleIORank, MPI_COMM_WORLD);
68
69     const Image img = { .pixels=in_image, width, height, max };
70     int subset_size = img_partition_pixels(img.width*img.height, cluster_size);
71     int subset_start = img_partition_head(subset_size, rank);
72
73
74 /*
75 * The "fuzzy" transformation
76 *
77 */
78 Pixel* out_image_fuzzy = (Pixel *) malloc(PPM_MAX_PIXELS * sizeof (Pixel));
79 assert(out_image_fuzzy != 0);
80
81 t0 = MPI_Wtime();
82 int xradius = 5, yradius = 5;
83 averaging_filter_subset(img, subset_start, subset_size, xradius, yradius,
84                         out_image_fuzzy);
85 MPI_Gather(out_image_fuzzy+subset_start, subset_size, MPI_PIXEL,
86             out_image_fuzzy, subset_size, MPI_PIXEL,
87             ConsoleIORank, MPI_COMM_WORLD);
88 tf = MPI_Wtime();
89
90 if (rank == ConsoleIORank) {
91     char* fuzzy_fname = append_to_basename(image_filename,
92                                            "-parallel-fuzzy.ppm");
93     err = write_ppm(fuzzy_fname, width, height, (char *) out_image_fuzzy);
94     if (err != 0) MPI_Abort(MPI_COMM_WORLD, err);
95
96     printf("[%d]: Averaging Filter: Neighborhood radius |x y|: %d %d\n",
97            rank, xradius, yradius);
98
99     t_delta_seconds = (tf - t0) * seconds_per_tick;
100    printf("[%d]: Averaging Filter: %d node parallel computation time: ",
101           rank, cluster_size);
102    printf("%g seconds\n", t_delta_seconds);
103 }
104
105
106 /*
107 * The threshold transformation
108 *
109 */
110 GPixel* out_image_thresh = (GPixel *) malloc(PPM_MAX_PIXELS * sizeof (GPixel));
111 assert(out_image_thresh != 0);

```

```

112     t0 = MPI_Wtime();
113     double local_intensity=sum_subset_intensity(img, subset_start, subset_size);
114     double global_threshold;
115     MPI_Allreduce(&local_intensity, &global_threshold, 1, MPI_DOUBLE, MPI_SUM,
116                   MPI_COMM_WORLD);
117     global_threshold /= img.width * img.height;
118     threshold_filter_subset(img, subset_start, subset_size, global_threshold,
119                             out_image_thresh);
120
121     MPI_Gather(out_image_thresh + subset_start, subset_size, MPI_UNSIGNED_CHAR,
122                out_image_thresh, subset_size, MPI_UNSIGNED_CHAR,
123                ConsoleIORank, MPI_COMM_WORLD);
124     tf = MPI_Wtime();
125
126     if (rank == ConsoleIORank) {
127         char* threshold_fname = append_to_basename(image_filename,
128                                              "-parallel-threshold.pgm");
129         err = write_pgm(threshold_fname, width, height,
130                         (char *) out_image_thresh);
131         if (err != 0) exit(err);
132
133         printf("[%d]: Threshold Filter: Overall average intensity is: %f\n",
134                rank, global_threshold);
135         t_delta_seconds = (tf - t0) * seconds_per_tick;
136         printf("[%d]: Threshold Filter: %d node parallel computation time: ",
137                rank, cluster_size);
138         printf("%g seconds\n", t_delta_seconds);
139         fflush(stdout);
140     }
141
142     MPI_Finalize();
143     return EXIT_SUCCESS;
144 }
145
146 void make_ppm_pixel_dt(MPI_Datatype* MPI_PIXEL) {
147     int lengths[3] = {1,1,1};
148     const MPI_Aint displacements[3] = {
149         0,
150         sizeof (unsigned char),
151         2 * sizeof (unsigned char)
152     };
153     MPI_Datatype types[3] = {
154         MPI_UNSIGNED_CHAR, MPI_UNSIGNED_CHAR, MPI_UNSIGNED_CHAR
155     };
156     MPI_Type_create_struct(3, lengths, displacements, types, MPI_PIXEL);
157     MPI_Type_commit(MPI_PIXEL);
158 }
159
160 /* example:
161 *
162 *  * append_to_basename("img.ppm", "-fuzzy.pgm")
163 *  *      ->"img-fuzzy.pgm"
164 */
165 char* append_to_basename(const char* filename, const char* appendix) {
166     size_t newlen = strlen(filename) + strlen(appendix);
167

```

```
168     char* newfile = (char *) malloc(sizeof (char) * newlen);
169     char* ft_extension = strrchr(filename, '.');
170     size_t basename_len = strlen(filename) - strlen(ft_extension);
171     strncpy(newfile, filename, basename_len); // without the file extension
172     strcat(newfile, appendix);
173     //strcat(newfile, ft_extension);
174     //printf("outfile name: %s\n", newfile);
175
176     return newfile;
177 }
```

**src/filter\_ppm.h**

```
1 /* filter_ppm.h
2 * -----
3 * Authors: Darwin Jacob Groskleg
4 *
5 * Purpose: extract filtering methods from both programs to make them more
6 *           testable by comparison.
7 *           Extension of ppm image struct.
8 *
9 * Filters: take image data & image metadata (width, height),
10 *           mutating the output img pointer data.
11 */
12 #ifndef FILTER_PPM_H_INCLUDED
13 #define FILTER_PPM_H_INCLUDED
14
15 #include "ppm.h"
16 #include <math.h>
17
18 void averaging_filter(const Image img, int x_radius, int y_radius,
19     Pixel* extant_img);
20 void averaging_filter_subset(const Image img, int start_pos, int count,
21     int x_radius, int y_radius,
22     Pixel* extant_img);
23
24 double threshold_filter(const Image img, GPixel* extant_img);
25 double sum_subset_intensity(const Image img, int start_pos, int count);
26 void threshold_filter_subset(const Image img, int start_pos, int count,
27     double threshold, GPixel* extant_img);
28
29
30 /* Helpers for Subset Filtering
31 * they're only helpers so go ahead and show the implementation for now.
32 */
33
34 /* how many pixels in a partition? */
35 inline int img_partition_pixels(int pixel_area, int partitions) {
36     return ceil(pixel_area / (float) partitions);
37 }
38
39 /* partition_size == pixels in a partition */
40 inline int img_partition_head(int partition_size, int kth_subset) {
41     return partition_size * kth_subset;
42 }
43
44
45#endif // FILTER_PPM_H_INCLUDED
```

## src/filter\_ppm.c

```

1  /* filter_ppm.c
2  * -----
3  * Authors: Darwin Jacob Groskleg
4  *
5  * Purpose: should be agnostic to using MPI.
6  */
7  #include "filter_ppm.h"
8
9  #include <assert.h>
10 #include <stddef.h>
11
12
13 #define min(a,b) \
14     ({ __typeof__ (a) _a = (a); \
15         __typeof__ (b) _b = (b); \
16         _a < _b ? _a : _b; })
17
18 /* Provide "external definitions" of the helper functions,
19 * required for proper linking when using C99 inline keyword.
20 */
21 extern inline int img_partition_pixels(int pixel_area, int partitions);
22 extern inline int img_partition_head(int partition_size, int kth_subset);
23
24
25 /* Averaging / Fuzzy Filter
26 *
27 * Simplest method for parallel:
28 * - (broadcast) send all data to all procs
29 * - partition the resulting pixel stream by process number
30 * - collect / gather by root, which writes to file
31 *
32 *
33 * Implementation Suggestions:
34 *
35 * The value for a pixel  $(x; y)$  in the output image is the sum of all the pixels
36 * in a rectangle in the input image centered around  $(x; y)$  divided by the
37 * number of pixels in the rectangle.
38 *
39 * In other words, the pixel in the output image is an average of the
40 * rectangle-shaped neighborhood of corresponding pixel in the input image.
41 *
42 * Your implementation should work with different sizes of the rectangle.
43 * Therefore, it is recommended that the function implementing this
44 * transformation take parameters specifying the x-radius and y-radius of the
45 * rectangle (see Figure 1). It is allowed to use integer limits for the
46 * rectangle to avoid the problem of interpolation.
47 */
48 void averaging_filter_subset(const Image img, int start_pos, int count,
49     int x_radius, int y_radius,
50     Pixel* extant_img)
51 {
52     assert(extant_img != 0);      // memory allocated?
53
54     // radius for averaging
55     int xr = x_radius;

```

```

56     int yr = y_radius;
57     assert(xr > 0 && yr > 0);
58
59     int sentinel = min(start_pos + count, img.width * img.height);
60     // k indexing = k'th element is a_ij
61     for (int x,y,k=start_pos; k < sentinel; k++) {
62         // row index = i
63         y = (int) k / img.width;      // integer division
64         // column index = j
65         x = k % img.width;
66
67         int r, g, b, n;
68         n = 0; r = 0; g = 0; b = 0;
69         for (int xas = x - xr; xas < x + xr; xas++) {
70             for (int yas = y - yr; yas < y + yr; yas++) {
71                 Pixel *pix;
72
73                 if ( (xas < 0) || (xas >= img.width)
74                     || (yas < 0) || (yas >= img.height))
75                 {
76                     pix = NULL;
77                 } else {
78                     pix = img.pixels + yas*img.width + xas;
79                 }
80
81                 if (pix != NULL) {
82                     r += pix->red;
83                     g += pix->green;
84                     b += pix->blue;
85                     n++;
86                 }
87             }
88         }
89
90         if (n > 0) {
91             extant_img[k].red   = r/n;
92             extant_img[k].green = g/n;
93             extant_img[k].blue  = b/n;
94         }
95     }
96 }
97
98 /* Averaging Filter
99 */
100 * Purpose: makes the image "fuzzy".
101 */
102 void averaging_filter(const Image img, int x_radius, int y_radius,
103                         Pixel* extant_img)
104 {
105     assert(extant_img != 0);
106
107     // radius for averaging
108     int xr = x_radius;
109     int yr = y_radius;
110     assert(xr > 0 && yr > 0);
111 }
```

```

112     for (int y = 0; y < img.height; y++) {
113         for (int x = 0; x < img.width; x++) {
114             int r, g, b, n;
115             n = 0; r = 0; g = 0; b = 0;
116             for (int xas = x - xr; xas < x + xr; xas++) {
117                 for (int yas = y - yr; yas < y + yr; yas++) {
118                     Pixel * pix;
119                     if ((xas < 0) || (xas >= img.width)
120                         || (yas < 0) || (yas >= img.height))
121                     {
122                         pix = NULL;
123                     } else {
124                         pix = img.pixels + yas*img.width + xas;
125                     }
126
127                     if (pix != NULL) {
128                         r += pix->red;
129                         g += pix->green;
130                         b += pix->blue;
131                         n++;
132                     }
133                 }
134             }
135
136             if (n > 0) {
137                 extant_img->red = r/n;
138                 extant_img->green = g/n;
139                 extant_img->blue = b/n;
140             }
141             extant_img++;
142         }
143     }
144 }
145
146 double average_intensity(const Image img) {
147     Pixel* pix;
148     double intensity = 0;
149     for (int y=0; y<img.height; y++) {
150         for (int x=0; x<img.width; x++) {
151             pix = img.pixels + y*img.width + x;
152             intensity += ( (1.0 * (double) pix->red
153                             + 1.0 * (double) pix->green
154                             + 1.0 * (double) pix->blue
155                             )/ 3.0);
156         }
157     }
158
159     intensity /= img.width * img.height;
160     return intensity;
161 }
162
163 /* Threshold Filter
164 *
165 * Purpose: creates a binary-pixel image.
166 *
167 * The filter computes the average intensity of the whole input image and use

```

```

168 * this value to threshold the image. The result is an image containing only
169 * black and white pixels. The white for those pixels in the input image that
170 * are lighter than the threshold and the black for those pixels in the input
171 * image that are darker than the threshold.
172 */
173 double threshold_filter(const Image img, GPixel* extant_img) {
174     assert(extant_img != 0);
175
176     double intensity = average_intensity(img);
177
178     Pixel* pix;
179     for (int y=0; y<img.height; y++) {
180         for (int x=0; x<img.width; x++) {
181             pix = img.pixels + y*img.width + x;
182             double pix_intensity = 0;
183             pix_intensity = ( 1.0 * (double) pix->red
184                             + 1.0 * (double) pix->green
185                             + 1.0 * (double) pix->blue
186                             ) / 3;
187             if ( pix_intensity > intensity) {
188                 extant_img[y*img.width + x] = 255;
189             } else {
190                 extant_img[y*img.width + x] = 0;
191             }
192         }
193     }
194     return intensity;
195 }
196
197 /* for use in a threshold filter
198 */
199 double sum_subset_intensity(const Image img, int start_pos, int count) {
200     Pixel* pix;
201     double intensity = 0;
202     int sentinel = min(start_pos + count, img.width * img.height);
203     // k indexing = k'th element is a_ij
204     for (int x,y,k=start_pos; k < sentinel; k++) {
205         // row index = i
206         y = (int) k / img.width;      // integer division
207         // column index = j
208         x = k % img.width;
209         pix = img.pixels + y*img.width + x;
210         intensity += ((1.0 * (double) pix->red
211                         + 1.0 * (double) pix->green
212                         + 1.0 * (double) pix->blue
213                         )/ 3.0);
214     }
215
216     return intensity;
217 }
218
219 /* Threshold Filter
220 *
221 * Simplest method:
222 * - scatter initial pixel stream, partitioned, to procs
223 * - calculate avg in each

```

```
224 * - Reduce and scatter the result, the average intensity
225 * - each proc applies threshold to values
226 * - collect / gather by root, which writes to file
227 */
228 void threshold_filter_subset(const Image img, int start_pos, int count,
229     double threshold,
230     GPixel* extant_img)
231 {
232     assert(extant_img != 0);
233     assert(threshold >= 0 && threshold <= 255);
234
235     int sentinel = min(start_pos + count, img.width * img.height);
236     // k indexing = k'th element is a_ij
237     for (int x,y,k=start_pos; k < sentinel; k++) {
238         // row index = i
239         y = (int) k / img.width;      // integer division
240         // column index = j
241         x = k % img.width;
242
243         Pixel *pix = img.pixels + y*img.width + x;
244         double pix_intensity = 0;
245         pix_intensity = ( 1.0 * (double) pix->red
246                         + 1.0 * (double) pix->green
247                         + 1.0 * (double) pix->blue
248                         ) / 3;
249         if ( pix_intensity > threshold) {
250             extant_img[y*img.width + x] = 255;
251         } else {
252             extant_img[y*img.width + x] = 0;
253         }
254     }
255 }
```

**src/ppm.h**

```
1  /* ppm.h
2   * -----
3   * Authors: Darwin Jacob Groskleg, Dr Laurence T. Yang
4   * Purpose: interface writing and reading Portable Pix Map (ppm) image format
5   *           files, the datatypes to hold it in.
6   */
7  #ifndef PPM_H_INCLUDED
8  #define PPM_H_INCLUDED
9
10 #define PPM_MAX_PIXELS (1000*1000)
11
12
13 /* image.h : MPI agnostic */
14 typedef struct Pixel {
15     unsigned char red, green, blue;
16 } Pixel;
17
18 /* Gray Pixel: 0 (black) – 255 (white) */
19 typedef unsigned char GPixel;
20
21 typedef struct Image {
22     /* PPM Body */
23     Pixel* pixels;
24     /* PPM Header */
25     int width;
26     int height;
27     int max_color;
28 } Image;
29
30
31 /* ppm.h : Single Threaded Only! */
32 int read_ppm(const char * fname, int* xpix, int* ypix, int* max, char* data);
33 int write_ppm(const char * fname, int xpix, int ypix, char* data);
34 int write_pgm(const char * fname, int xpix, int ypix, char* data);
35
36#endif // PPM_H_INCLUDED
```

## src/ppm.c

```

1  /* ppm.c
2  * -----
3  * Authors: Darwin Jacob Groskleg, Dr Laurence T. Yang
4  *
5  * Purpose: implementation for writing and reading Portable Pix Map (ppm)
6  *           image format files.
7  */
8 #include "ppm.h"
9
10 #include <errno.h>
11 #include <assert.h>
12 #include <stdio.h>
13 #include <string.h>
14
15 /*
16  * Input arguments:
17  *     *fname : filename to open
18  *
19  * Output arguments:
20  *     *xpix :
21  *     *ypix :
22  *     *max :
23  *     data[] :
24 */
25 int read_ppm(const char* fname, int* xpix, int* ypix, int* max, char* data) {
26     char ftype[40];
27     char ctype[40] = "P6"; // binary data, more efficient than P3
28     char line[80];
29
30     FILE * fp;
31     errno = 0;
32
33     if (fname == NULL) fname = "\0";
34     fp = fopen(fname, "r");
35     if (fp == NULL) {
36         fprintf(stderr,
37                 "read_ppm failed to open %s: %s\n", fname, strerror(errno));
38         return 1;
39     }
40
41     fgets(line, 80, fp); // first 80 char or until newline
42     sscanf(line, "%s", ftype); // assign string from `line` to ftype[]
43
44     while ( fgets(line, 80, fp) && (line[0] == '#') ); // skip comments
45
46     // grab the pixel width/height numbers from `line`
47     sscanf(line, "%d%d", xpix, ypix);
48
49     // grab the next number from the stream,
50     // the maximum value on the colour range
51     fscanf(fp, "%d\n", max);
52
53 /*
54     fprintf(stderr, "typ= %s xmax= %d ymax= %d intensity max= %d\n",
55             ftype, *xpix, *ypix, *max);

```

```
56     fflush(stderr);
57 */
58
59     assert(*xpix * *ypix <= PPM_MAX_PIXELS);
60
61     if (strncmp(ftype, ctype, 2) == 0) {
62         size_t pixel_channels = *xpix * *ypix * 3; // pixels x 3
63         size_t position = fread(data, sizeof(char), pixel_channels, fp);
64         if (position != pixel_channels) {
65             perror("Read failed");
66             fclose(fp);
67             return 2;
68         }
69     } else {
70         fprintf(stderr, "Wrong file format: %s\n", ftype);
71     }
72
73     if (fclose(fp) == EOF) {
74         perror("Close failed");
75         return 3;
76     }
77
78     return 0;
79 }
80
81
82 int write_ppm(const char * fname, int xpix, int ypix, char * data) {
83     FILE * fp;
84     errno = 0;
85
86     if (fname == NULL) fname = "\0";
87     fp = fopen(fname, "w");
88     if (fp == NULL) {
89         fprintf(stderr,
90                 "write_ppm failed to open %s: %s\n", fname, strerror(errno));
91         return 1;
92     }
93
94     fprintf(fp, "P6\n");
95     fprintf(fp, "%d %d 255\n", xpix, ypix);
96     size_t pixel_channels = xpix * ypix * 3; // pixels x 3
97     if (fwrite(data, sizeof(char), pixel_channels, fp) != pixel_channels) {
98         perror("Write failed");
99         fclose(fp);
100        return 2;
101    }
102    if (fclose(fp) == EOF) {
103        perror("Close failed");
104        return 3;
105    }
106    return 0;
107 }
108
109
110 int write_pgm(const char * fname, int xpix, int ypix, char * data) {
111     FILE * fp;
```

```
112     errno = 0;
113
114     if (fname == NULL) fname = "\0";
115     fp = fopen(fname, "w");
116     if (fp == NULL) {
117         fprintf(stderr, "write_ppm failed to open %s: %s\n", fname,
118                 strerror(errno));
119         return 1;
120     }
121
122     fprintf(fp, "P5\n");
123     fprintf(fp, "%d %d 255\n", xpix, ypix);
124     size_t pixels = xpix * ypix;
125     if (fwrite(data, sizeof(char), pixels, fp) != pixels) {
126         perror("Write failed");
127         fclose(fp);
128         return 2;
129     }
130     if (fclose(fp) == EOF) {
131         perror("Close failed");
132         return 3;
133     }
134     return 0;
135 }
```

**src/transform-serial.c**

```

1  /* transform-serial.c
2  * -----
3  * Derived from: serial.c (L.T. Yang)
4  * Authors: Darwin Jacob Groskleg, Dr Laurence T. Yang
5  *
6  * Usage:
7  *     transform-serial [input.ppm]
8  * Output:
9  * - file:      input-fuzzy.ppm
10 * - file:      input-threshold.pgm
11 * - stdout:    average intensity calculated for the fuzzy filter.
12 */
13 #include <stdio.h>
14 #include <errno.h>
15 #include <assert.h>
16 #include <stdlib.h>
17 #include <stddef.h>
18 #include <string.h>
19
20 #include "ppm.h"
21 #include "filter_ppm.h"
22
23 char* append_to_basename(const char* filename, const char* appendix);
24
25 int main(int argc, char *argv[]) {
26     const char* default_image_filename = "./im1.ppm";
27     char* image_filename;
28     if (argc > 1) {
29         image_filename = (char *) malloc( sizeof (char) * strlen(argv[1]));
30         strcpy(image_filename, argv[1]);
31     } else {
32         size_t default_len = strlen(default_image_filename);
33         image_filename = (char *) malloc( sizeof (char) * default_len);
34         strcpy(image_filename, default_image_filename);
35     }
36     printf("infile name: %s\n", image_filename);
37
38     Pixel* in_image = (Pixel *) malloc(PPM_MAX_PIXELS * sizeof (Pixel));
39     assert(in_image != 0);
40
41     int err = 0;
42
43     int width, height, max;
44     err = read_ppm(image_filename, &width, &height, &max, (char *) in_image);
45     if (err != 0) exit(err);
46     const Image img = { .pixels=in_image, width, height, max };
47
48 /*
49 * The "fuzzy" transformation
50 *
51 */
52
53     Pixel* out_image_fuzzy = (Pixel *) malloc(PPM_MAX_PIXELS * sizeof (Pixel));
54     assert(out_image_fuzzy != 0);
55

```

```
56     int xradius = 5, yradius = 5;
57     printf("Neighborhood pixel radius |x y|: %d %d\n", xradius, yradius);
58     averaging_filter(img, xradius, yradius, out_image_fuzzy);
59
60     char* fuzzy_fname = append_to_basename(image_filename, "-serial-fuzzy.ppm");
61     err = write_ppm(fuzzy_fname, width, height, (char *) out_image_fuzzy);
62     if (err != 0) exit(err);
63
64 /*
65 * The thresholding
66 */
67
68 GPixel* out_image_thresh = (GPixel *) malloc(PPM_MAX_PIXELS * sizeof (GPixel));
69 assert(out_image_thresh != 0);
70
71 double intensity = threshold_filter(img, out_image_thresh);
72 printf("Average intensity is: %f\n", intensity);
73 fflush(stdout);
74
75 char* threshold_fname = append_to_basename(image_filename,
76                                              "-serial-threshold.pgm");
77 err = write_pgm(threshold_fname, width, height, (char *) out_image_thresh);
78 if (err != 0) exit(err);
79
80
81 return 0;
82 }
83
84 /* example:
85 *
86 * append_to_basename("img.ppm", "-fuzzy.pgm")
87 *      ->"img-fuzzy.pgm"
88 */
89
90 char* append_to_basename(const char* filename, const char* appendix) {
91     size_t newlen = strlen(filename) + strlen(appendix);
92     char* newfile = (char *) malloc(sizeof (char) * newlen);
93     char* ft_extension = strrchr(filename, '.');
94     size_t basename_len = strlen(filename) - strlen(ft_extension);
95     strncpy(newfile, filename, basename_len); // without the file extension
96     strncat(newfile, appendix);
97     //strcat(newfile, ft_extension);
98     //printf("outfile name: %s\n", newfile);
99
100    return newfile;
101 }
```

**Makefile**

```

1 include ../mpi.mk
2 include ../clean.mk
3
4 vpath %.h      ./src
5 vpath %.c      ./src
6
7 clean: clean-ppm-images
8
9 all: transform-serial transform-parall test_filter
10
11 ppm.o :          ppm.c ppm.h
12 filter_ppm.o :   filter_ppm.c filter_ppm.h ppm.h
13
14 transform-serial.o : transform-serial.c ppm.h filter_ppm.h
15 transform-serial : transform-serial.o ppm.o filter_ppm.o
16
17 transform-parall.o : transform-parall.c ppm.h filter_ppm.h
18 transform-parall : transform-parall.o ppm.o filter_ppm.o
19
20 test_filter.o :   test_filter.c ppm.h filter_ppm.h
21 test_filter :     test_filter.o ppm.o filter_ppm.o
22
23 .PHONY:
24 test-filter: test_filter
25     ./$<
26
27 .PHONY: test-serial
28 test-serial: transform-serial
29     ./$<
30
31 .PHONY: test-parall
32 test-parall: transform-parall
33     $(run-mpi)
34
35 .PHONY: test-parall-all
36 test-parall-all: transform-parall
37     $(run-mpi) im1.ppm
38     $(run-mpi) im2.ppm
39     $(run-mpi) im3.ppm
40
41 # Run 4 times, drop the first for disk caching outliers
42 .PHONY: perf-q1
43 perf-q1: transform-parall transform-serial
44     time ./transform-serial im1.ppm
45     time $(MPIRUN) -np 1 ./transform-parall im1.ppm
46     @echo -- Reject First, Keep Below -----
47     time ./transform-serial im1.ppm
48     time $(MPIRUN) -np 1 ./transform-parall im1.ppm
49     time ./transform-serial im1.ppm
50     time $(MPIRUN) -np 1 ./transform-parall im1.ppm
51     time ./transform-serial im1.ppm
52     time $(MPIRUN) -np 1 ./transform-parall im1.ppm

```

```

53 .PHONY: perf-q3
54 perf-q3: transform-parall
55     time $(MPIRUN) -np 1 ./transform-parall im1.ppm
56     @echo -- Reject First, Keep Below -----
57     time $(MPIRUN) -np 1 ./transform-parall im1.ppm
58     time $(MPIRUN) -np 2 ./transform-parall im1.ppm
59     time $(MPIRUN) -np 4 ./transform-parall im1.ppm
60
61
62 # invoked by :clean
63 .PHONY: clean-pmm-images
64 clean-pmm-images:
65     rm im?-*.*ppm im?-*.*pgm
66
67 # Merge all 3 to a single before and after image
68 # requires ImageMagick
69 assets/%-merged.png : %.ppm %-parall-fuzzy.ppm %-parall-threshold.pgm
70     magick montage -geometry 100% $^ $@
71
72 # old ImageMagick command
73 # convert $^ +append $@
74 # from netpbm-package
75 # pnmcat -lr <(pngtopnm 1.png) <(pngtopnm 2.png) | pnmtopng > all.png
76
77 project_label:=Lab5-6
78 ordered_docs := questions.md \
79                 src/transform-parall.c \
80                 src/filter_ppm.h   src/filter_ppm.c \
81                 src/ppm.h         src/ppm.c \
82                 src/transform-serial.c \
83                 Makefile
84
85 vpath %.png ./assets
86 document_assets:= im1-merged.png \
87                   im2-merged.png \
88                   im3-merged.png \
89                   Screenshot-serial.png \
90                   Screenshot-parall-all.png
91
92 TITLE:=$(shell cat title.txt)
93 include ../publish.mk

```