CSCI 455: Lab #8 — Advanced MPI: Group/Communicator Management and Virtual Topologies

Darwin Jacob Groskleg

Winter 2020

Contents

Part 1: Group/Communicator Management

MPI Split Method

```
darwingroskleg@starbuck ~/Dropbox/Documents/Terms/2020-01 - Winter/CSCI455/Lab8-Adv MPI Communicator Mgmt
[make p1-1
mpicc -std=c99 -Wall -Wextra -g -D_GLIBCXX_DEBUG -O0  -lmpi  group-split.c   -o group-split
Platform: Darwin (4 cpu cores recognized)
PMIX_MCA_gds=hash mpirun --host localhost --mca btl_vader_backing_directory /tmp --mca btl ^tcp --oversubscribe
-np 4 ./group-split
Old_proc [0] has new rank [0], received value 0
Old_proc [1] has new rank [0], received value 1
Old_proc [2] has new rank [1], received value 0
Old_proc [3] has new rank [1], received value 1
darwingroskleg@starbuck ~/Dropbox/Documents/Terms/2020-01 - Winter/CSCI455/Lab8-Adv MPI Communicator Mgmt
```

Figure 1: Console output: group-split.c

group-split.c

```c
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[]) {
    MPI_Init(&argc, &argv);
    int numprocs;
    MPI_Comm_size(MPI_COMM_WORLD, &numprocs);
    int world_rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);

    /* Color of local process:
     *  decides to which communicator the process will
     *  belong after the split.
     */
    int color = world_rank % 2;

    // Communicator split is a remote opration?
    MPI_Comm New_Comm;
    MPI_Comm_split(MPI_COMM_WORLD, color, world_rank, &New_Comm);
    int new_rank;
    MPI_Comm_rank(New_Comm, &new_rank);
    int new_nodes;
    MPI_Comm_size(New_Comm, &new_nodes);

    int broad_val;
    if(new_rank == 0)
        broad_val = color;
    MPI_Bcast(&broad_val, 1, MPI_INT, 0, New_Comm);

    printf("Old_proc [%d] has new rank [%d], received value %d\n",
              world_rank,          new_rank,           broad_val);

    MPI_Comm_free(&New_Comm);
    MPI_Finalize();
    return 0;
}
```

MPI Include Method

```
ubuntu@ip-172-31-43-252:~/csci455/Lab8-Adv_MPI_Communicator_Mgmt$ make p1-2
mpicc -std=c99 -Wall -Wextra -g -D_GLIBCXX_DEBUG -O0    group-incl.c  -lm -o group-incl
Platform: Linux (96 cpu cores recognized)
MPIRUN group-incl with 8 node processes:
rank = 0  newrank = 0  recvbuf= 6
rank = 1  newrank = 1  recvbuf= 6
rank = 2  newrank = 2  recvbuf= 6
rank = 3  newrank = 3  recvbuf= 6
rank = 4  newrank = 0  recvbuf= 22
rank = 5  newrank = 1  recvbuf= 22
rank = 6  newrank = 2  recvbuf= 22
rank = 7  newrank = 3  recvbuf= 22
ubuntu@ip-172-31-43-252:~/csci455/Lab8-Adv_MPI_Communicator_Mgmt$
```

Figure 2: Console output: group-incl.c running on a 96 node compute cluster.

group-incl.c

```c
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>

#define NPROCS 8

int main(int argc, char *argv[])  {
    MPI_Init(&argc,&argv);
    int world_rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);
    int numtasks;
    MPI_Comm_size(MPI_COMM_WORLD, &numtasks);

    if (numtasks != NPROCS) {
        printf("Must specify MP_PROCS= %d. Terminating.\n",NPROCS);
        MPI_Finalize();
        exit(0);
    }

    // extract the original group handle
    MPI_Group  orig_group;
    MPI_Comm_group(MPI_COMM_WORLD, &orig_group);
    MPI_Group  new_group;

    //  divide tasks into two distinct groups based upon world_rank
    //  Group Contruction: happens locally.
    if (world_rank < NPROCS/2) {
        const int ranks1[4]={0,1,2,3};
        MPI_Group_incl(orig_group, 4, ranks1, &new_group);
    }
    else {
        const int ranks2[4]={4,5,6,7};
        MPI_Group_incl(orig_group, 4, ranks2, &new_group);
    }

    // create new new communicator and then perform collective communications
    // Communicator Construction: happens remotely/collectively.
    MPI_Comm   new_comm;
    MPI_Comm_create(MPI_COMM_WORLD, new_group, &new_comm);

    int sendbuf = world_rank;
    int recvbuf;
    MPI_Allreduce(&sendbuf, &recvbuf, 1, MPI_INT, MPI_SUM, new_comm);

    // get rank in new group
    int new_rank;
    MPI_Group_rank(new_group, &new_rank);


    printf("rank = %d  newrank = %d  recvbuf= %d\n",
           world_rank,     new_rank,     recvbuf);


    MPI_Group_free(&orig_group);
    MPI_Group_free(&new_group);
```

```
56      MPI_Comm_free(&new_comm);
57      MPI_Finalize();
58      return 0;
59  }
```

Part 2: Virtual Topologies

MPI Cartesian Map with 1 Dimension

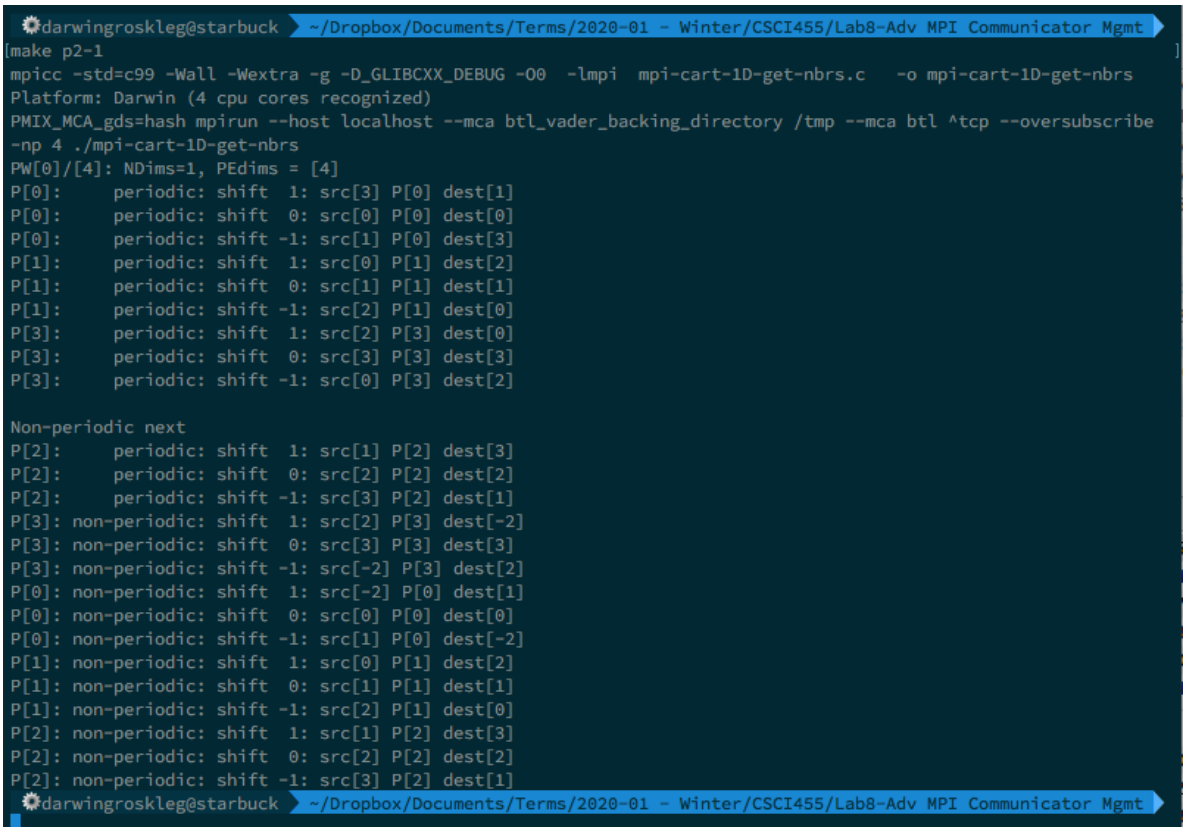

Figure 3: Console output: mpi-cart-1D-get-nbrs.c

mpi-cart-1D-get-nbrs.c

```c
/* mpi-cart-1D-get-nbrs.c
 * ----------------------
 * 1 Dimensional Cartesian Virtual Topology
 *  finds the neighbors in a cartesian communicator
 */
#include <mpi.h>
#include <stdlib.h>
#include <stdio.h>

typedef enum { false, true } bool;

int main( int argc, char *argv[] ) {
    MPI_Init( &argc, &argv );
    int cluster_size;
    MPI_Comm_size( MPI_COMM_WORLD, &cluster_size );
    int rank;
    MPI_Comm_rank( MPI_COMM_WORLD, &rank );

    const int ndims = 1;
    int dims[ndims];
    /* processor dimensions */
    dims[0] = cluster_size;
    /* create Cartesian topology for processes */
    //      nnodes, ndims, dims[]
    MPI_Dims_create(cluster_size, ndims, dims);

    if(rank == 0)
        printf("PW[%d]/[%d]: NDims=%d, PEdims = [%d]\n",
                rank, cluster_size, ndims,  dims[0]);

    int periods[ndims];
    int source, dest;
    bool reorder = true;
    MPI_Comm comm1D;

    /*****************************************************************/
    /* Create periodic shift */
    /*****************************************************************/
    /* periodic shift is true. */
    periods[0] = true;
    /* create Cartesian mapping */
    MPI_Cart_create(MPI_COMM_WORLD, ndims, dims, periods, reorder, &comm1D);

    //int errs = 0;
    MPI_Cart_shift(comm1D, dims[0], 1, &source, &dest);
    printf( "P[%d]:     periodic: shift  1: src[%d] P[%d] dest[%d]\n",
             rank,                      source, rank,    dest );
    fflush(stdout);

    MPI_Cart_shift(comm1D, dims[0], 0, &source, &dest);
    printf( "P[%d]:     periodic: shift  0: src[%d] P[%d] dest[%d]\n",
             rank,                      source, rank,    dest );
    fflush(stdout);

    MPI_Cart_shift(comm1D, dims[0], -1, &source, &dest);
```

```
56      printf( "P[%d]:     periodic: shift -1: src[%d] P[%d] dest[%d]\n",
57              rank,                       source, rank,    dest );
58      fflush(stdout);
59      MPI_Comm_free( &comm1D );
60
61      /*************************************************************/
62      /* Create non-periodic shift */
63      /*************************************************************/
64      if (rank == 0)
65          printf("\nNon-periodic next\n");
66      /* periodic shift is false. */
67      periods[0] = false;
68      MPI_Cart_create(MPI_COMM_WORLD, ndims, dims, periods, reorder, &comm1D);
69
70      MPI_Cart_shift(comm1D, dims[0], 1, &source, &dest);
71      printf( "P[%d]: non-periodic: shift  1: src[%d] P[%d] dest[%d]\n",
72              rank,                       source, rank,    dest );
73      fflush(stdout);
74      MPI_Cart_shift(comm1D, dims[0], 0, &source, &dest);
75      printf( "P[%d]: non-periodic: shift  0: src[%d] P[%d] dest[%d]\n",
76              rank,                       source, rank,    dest );
77      fflush(stdout);
78      MPI_Cart_shift(comm1D, dims[0], -1, &source, &dest);
79      printf( "P[%d]: non-periodic: shift -1: src[%d] P[%d] dest[%d]\n",
80              rank,                       source, rank,    dest );
81      fflush(stdout);
82      MPI_Comm_free( &comm1D );
83
84      MPI_Finalize();
85      return 0;
86  }
```

MPI Cartesian Map with 2 Dimensions

```
ubuntu@ip-172-31-43-252:~/csci455/Lab8-Adv_MPI_Communicator_Mgmt$ make p2-2
mpicc -std=c99 -Wall -Wextra -g -D_GLIBCXX_DEBUG -O0    mpi-cart-2D-get-nbrs.c  -lm -o mpi-cart-2D-get-nbrs
Platform: Linux (96 cpu cores recognized)
MPIRUN mpi-cart-2D-get-nbrs with 16 node processes:
PW[0], CommSz[16]: PEdims = [4 x 4]
PW[ 1] Coord(0,1): SHIFT_DIM[0], Shift=1: nbr_lo[ 0] P[ 1] nbr_hi[ 2]
PW[ 1] Coord(0,1): SHIFT_DIM[1], Shift=1: nbr_lo[ 0] P[ 1] nbr_hi[ 2]
PW[ 2] Coord(0,2): SHIFT_DIM[0], Shift=1: nbr_lo[ 1] P[ 2] nbr_hi[ 3]
PW[ 2] Coord(0,2): SHIFT_DIM[1], Shift=1: nbr_lo[ 1] P[ 2] nbr_hi[ 3]
PW[ 3] Coord(0,3): SHIFT_DIM[0], Shift=1: nbr_lo[ 2] P[ 3] nbr_hi[ 2]
PW[ 3] Coord(0,3): SHIFT_DIM[1], Shift=1: nbr_lo[ 2] P[ 3] nbr_hi[ 2]
PW[ 4] Coord(1,0): SHIFT_DIM[0], Shift=1: nbr_lo[ 2] P[ 4] nbr_hi[ 5]
PW[ 4] Coord(1,0): SHIFT_DIM[1], Shift=1: nbr_lo[ 2] P[ 4] nbr_hi[ 5]
PW[ 5] Coord(1,1): SHIFT_DIM[0], Shift=1: nbr_lo[ 4] P[ 5] nbr_hi[ 6]
PW[ 5] Coord(1,1): SHIFT_DIM[1], Shift=1: nbr_lo[ 4] P[ 5] nbr_hi[ 6]
PW[ 6] Coord(1,2): SHIFT_DIM[0], Shift=1: nbr_lo[ 5] P[ 6] nbr_hi[ 7]
PW[ 6] Coord(1,2): SHIFT_DIM[1], Shift=1: nbr_lo[ 5] P[ 6] nbr_hi[ 7]
PW[ 7] Coord(1,3): SHIFT_DIM[0], Shift=1: nbr_lo[ 6] P[ 7] nbr_hi[ 2]
PW[ 7] Coord(1,3): SHIFT_DIM[1], Shift=1: nbr_lo[ 6] P[ 7] nbr_hi[ 2]
PW[ 8] Coord(2,0): SHIFT_DIM[0], Shift=1: nbr_lo[ 2] P[ 8] nbr_hi[ 9]
PW[ 8] Coord(2,0): SHIFT_DIM[1], Shift=1: nbr_lo[ 2] P[ 8] nbr_hi[ 9]
PW[ 9] Coord(2,1): SHIFT_DIM[0], Shift=1: nbr_lo[ 8] P[ 9] nbr_hi[10]
PW[ 9] Coord(2,1): SHIFT_DIM[1], Shift=1: nbr_lo[ 8] P[ 9] nbr_hi[10]
PW[10] Coord(2,2): SHIFT_DIM[0], Shift=1: nbr_lo[ 9] P[10] nbr_hi[11]
PW[10] Coord(2,2): SHIFT_DIM[1], Shift=1: nbr_lo[ 9] P[10] nbr_hi[11]
PW[11] Coord(2,3): SHIFT_DIM[0], Shift=1: nbr_lo[10] P[11] nbr_hi[ 2]
PW[11] Coord(2,3): SHIFT_DIM[1], Shift=1: nbr_lo[10] P[11] nbr_hi[ 2]
PW[12] Coord(3,0): SHIFT_DIM[0], Shift=1: nbr_lo[ 2] P[12] nbr_hi[13]
PW[12] Coord(3,0): SHIFT_DIM[1], Shift=1: nbr_lo[ 2] P[12] nbr_hi[13]
PW[13] Coord(3,1): SHIFT_DIM[0], Shift=1: nbr_lo[12] P[13] nbr_hi[14]
PW[13] Coord(3,1): SHIFT_DIM[1], Shift=1: nbr_lo[12] P[13] nbr_hi[14]
PW[14] Coord(3,2): SHIFT_DIM[0], Shift=1: nbr_lo[13] P[14] nbr_hi[15]
PW[14] Coord(3,2): SHIFT_DIM[1], Shift=1: nbr_lo[13] P[14] nbr_hi[15]
PW[15] Coord(3,3): SHIFT_DIM[0], Shift=1: nbr_lo[14] P[15] nbr_hi[ 2]
PW[15] Coord(3,3): SHIFT_DIM[1], Shift=1: nbr_lo[14] P[15] nbr_hi[ 2]
PW[ 0] Coord(0,0): SHIFT_DIM[0], Shift=1: nbr_lo[ 2] P[ 0] nbr_hi[ 1]
PW[ 0] Coord(0,0): SHIFT_DIM[1], Shift=1: nbr_lo[ 2] P[ 0] nbr_hi[ 1]
ubuntu@ip-172-31-43-252:~/csci455/Lab8-Adv_MPI_Communicator_Mgmt$
```

Figure 4: Console output: mpi-cart-2D-get-nbrs.c running on a 96 node compute cluster.

mpi-cart-2D-get-nbrs.c

```c
/* mpi-cart-2D-get-nbrs.c
 * ----------------------
 * 2 Dimensional Cartesian Virtual Topology
 *  finds the neighbors in a cartesian communicator
 */
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>
#include <math.h>

#define SHIFT_ROW 0
#define SHIFT_COL 1
#define DISP 1

typedef enum { false, true } bool;

int wrap_row(int  row_width, int col);
int wrap_col(int col_height, int row);

int main(int argc, char *argv[]) {
    //int errs;

    /* start up initial MPI environment */
    MPI_Init(&argc, &argv);
    int size;
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    int my_rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);

    int ndims=2;
    int dims[ndims];
    int nrows;
    int ncols;

    /* process command line arguments*/
    if (argc == 3) {
        nrows = atoi (argv[1]);
        ncols = atoi (argv[2]);
        dims[0] = nrows; /* number of rows */
        dims[1] = ncols; /* number of columns */
        if( (nrows*ncols) != size) {
            if( my_rank ==0)
                printf("ERROR: nrows*ncols) = %d * %d = %d != %d\n",
                        nrows, ncols, nrows*ncols, size);
            MPI_Finalize();
            exit(0);
        }
    }
    else {
        nrows = ncols = (int)sqrt(size);
        dims[0] = dims[1] = 0;
    }

    /*************************************************************/
    /* create cartesian topology for processes */
```

```
56      /*************************************************************/
57      MPI_Dims_create(size, ndims, dims);
58      if(my_rank == 0)
59          printf("PW[%d], CommSz[%d]: PEdims = [%d x %d]\n",
60                  my_rank,     size,       dims[0], dims[1]);
61
62      /* create cartesian mapping */
63      int periods[ndims];
64      periods[0] = periods[1] = 0; /* periodic shift is .false. */
65      int reorder = true;
66      MPI_Comm comm2D;
67      int ierr = 0;
68      ierr = MPI_Cart_create(MPI_COMM_WORLD, ndims, dims, periods, reorder,
69                             &comm2D);
70      if (ierr != 0)
71          printf("ERROR[%d] creating CART\n", ierr);
72
73      /* find my coordinates in the cartesian communicator group */
74      int coord[ndims];
75      MPI_Cart_coords(comm2D, my_rank, ndims, coord);
76
77      /* use my cartesian coordinates to find my rank in cartesian group*/
78      int my_cart_rank;
79      MPI_Cart_rank(comm2D, coord, &my_cart_rank);
80
81      //int source, dest;
82      /* get my neighbors; axis is coordinate dimension of shift */
83      /* axis=0 ==> shift along the rows: P[my_row-1]: P[me] : P[my_row+1] */
84      /* axis=1 ==> shift along the columns P[my_col-1]: P[me] : P[my_col+1] */
85      int nbr_i_lo, nbr_i_hi;
86      MPI_Cart_shift(comm2D, dims[0], DISP, &nbr_i_lo, &nbr_i_hi);
87      nbr_i_lo = wrap_row(ncols, nbr_i_lo);
88      nbr_i_hi = wrap_row(ncols, nbr_i_hi);
89      printf("PW[%2d] Coord(%d,%d): SHIFT_DIM[%d], Shift=%d: "
90                                  "nbr_lo[%2d] P[%2d] nbr_hi[%2d]\n",
91              my_rank, coord[0], coord[1], SHIFT_ROW,  DISP,
92                                       nbr_i_lo, my_rank,   nbr_i_hi);
93
94      int nbr_j_lo, nbr_j_hi;
95      MPI_Cart_shift(comm2D, dims[1], DISP, &nbr_j_lo, &nbr_j_hi);
96      nbr_j_lo = wrap_col(nrows, nbr_j_lo);
97      nbr_j_hi = wrap_col(nrows, nbr_j_hi);
98      printf("PW[%2d] Coord(%d,%d): SHIFT_DIM[%d], Shift=%d: "
99                                  "nbr_lo[%2d] P[%2d] nbr_hi[%2d]\n",
100             my_rank, coord[0], coord[1], SHIFT_COL,  DISP,
101                                      nbr_j_lo, my_rank,   nbr_j_hi);
102     fflush(stdout);
103
104     MPI_Comm_free( &comm2D );
105     MPI_Finalize();
106     return 0;
107 }
108
109 int wrap_dim(int dim_width, int idx) {
110     if (idx < 0)
111         return dim_width + idx;
```

```
112     return idx;
113 }
114 int wrap_row(int row_width, int col) {  return wrap_dim(row_width, col); }
115 int wrap_col(int col_height, int row) { return wrap_dim(col_height, row); }
```

Part 3: Simplified Matrix Multiplication

Using Cannon's algorithm.

```
ubuntu@ip-172-31-43-252:~/csci455/Lab8-Adv_MPI_Communicator_Mgmt$ make p3
mpicc -std=c99 -Wall -Wextra -g -D_GLIBCXX_DEBUG -O0    cannon.c  -lm -o cannon
Platform: Linux (96 cpu cores recognized)
MPIRUN cannon with 16 node processes:
A matrix:
0.000000        1.000000        2.000000        3.000000
4.000000        5.000000        6.000000        7.000000
8.000000        9.000000        10.000000       11.000000
12.000000       13.000000       14.000000       15.000000

B matrix:
1.000000        1.000000        1.000000        1.000000
1.000000        1.000000        1.000000        1.000000
1.000000        1.000000        1.000000        1.000000
1.000000        1.000000        1.000000        1.000000

A * B:
6.000000        6.000000        6.000000        6.000000
22.000000       22.000000       22.000000       22.000000
38.000000       38.000000       38.000000       38.000000
54.000000       54.000000       54.000000       54.000000
ubuntu@ip-172-31-43-252:~/csci455/Lab8-Adv_MPI_Communicator_Mgmt$
```

Figure 5: Console output: cannon.c running on a 96 node compute cluster.

cannon.c

```c
/* cannon.c
 * --------
 * Simplified Matrix-Matrix Multiplication
 *
 *  This code is based on Cannon algorithm for matrix matrix multiplication.
 *  The main assumption in Cannon is that both A and B matrix must be square
 *  matrix and number of processors must be equal to the no of elements in
 *  A matrix.
 */
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>

#define ndims 2 // 2 Dimension topology
#define SHIFT_ROW 1     // coord[1] is j
#define SHIFT_COL 0     // coord[0] is i

enum TaskRanks { Master = 0 };
// MxN Matrix: M rows, N columns
typedef struct MatrixSize {
  int m;    // rows    indexed by i
  int n;    // columns indexed by j
} mat_size_t;
mat_size_t get_matrices(float **L, float **R);

int main(int argc, char *argv[]) {
  // Initializing MPI
  MPI_Init(&argc, &argv);
  int size;
  MPI_Comm_size(MPI_COMM_WORLD, &size);
  int rank;
  MPI_Comm_rank(MPI_COMM_WORLD, &rank);

  // Read the data only if it is the root process (rank = 0)
  int   row, column;
  float *A = NULL;
  float *B = NULL;
  if (rank == Master) {
    mat_size_t msize = get_matrices(&A, &B);
    row    = msize.m;
    column = msize.n;
  }

  MPI_Barrier(MPI_COMM_WORLD);
  MPI_Bcast(&row,    1, MPI_INT, Master, MPI_COMM_WORLD);
  MPI_Bcast(&column, 1, MPI_INT, Master, MPI_COMM_WORLD);

  // set periodicity both vertical and horizontal movement
  //    periodic == true, wraps == true
  int periods[ndims] = {1, 1};
  int dims[ndims] = {row, column};
  int reorder = 1;  // true
  // Create Cartesian mapping of processes, a topological map
  MPI_Comm cart_comm;
  MPI_Cart_create(MPI_COMM_WORLD, ndims, dims, periods, reorder, &cart_comm);
```

```
56
57  // Sending/Assigning each A and B element to the individual processor
58  //    ASSUMES 1 element per process, enough processes
59  float a_ij = 0;
60  MPI_Scatter(A,    1, MPI_FLOAT,
61              &a_ij, 1, MPI_FLOAT,
62              Master, cart_comm);
63  float b_ij = 0;
64  MPI_Scatter(B,    1, MPI_FLOAT,
65              &b_ij, 1, MPI_FLOAT,
66              Master, cart_comm);
67  //printf("p[%d] a=%f, b=%f\n", rank, a_ij, b_ij);
68  MPI_Barrier(MPI_COMM_WORLD);
69
70  // 2 Dimension topology, so 2 coordinates
71  int coords[2];
72  // get the coordinates in the new Cartesian grid
73  MPI_Cart_coords(cart_comm, rank, ndims,   coords);
74  // get the new rank in Cartesian group using coords
75  int cart_rank;
76  MPI_Cart_rank(cart_comm, coords,  &cart_rank);
77  //printf("Coordinate of processor rank %d is [%d, %d], new rank is %d\n",
78  //                                  rank, coords[0], coords[1], cart_rank);
79
80  float c_ij = 0;
81  int msg_tag = 11;
82  // neighbor ranks
83  int right = 0, left = 0, down = 0, up = 0;
84  // Pumping along systolic array:
85  //    ASSUMES a square matrix
86  for (int ij = 0; ij < row; ij++) {
87    // get the shifted source and destination rank horizontally
88    MPI_Cart_shift(cart_comm, SHIFT_ROW, ij, &right, &left);
89    // get the shifted source and destination rank vertically
90    MPI_Cart_shift(cart_comm, SHIFT_COL, ij, &down, &up);
91    // send and receive using single buffer:
92    //      shift value from RIGHT coordinate to LEFT coordinate
93    MPI_Sendrecv_replace(&a_ij, 1, MPI_FLOAT,
94            left,  msg_tag,     // rank of dest   (send to   left)
95            right, msg_tag,     // rank of source (recv from right)
96            cart_comm, MPI_STATUS_IGNORE);
97    // send and receive using single buffer:
98    //      shift value from DOWN coordinate to UP coordinate
99    MPI_Sendrecv_replace(&b_ij, 1, MPI_FLOAT,
100           up,   msg_tag,      // rank of dest   (send up)
101           down, msg_tag,      // rank of source (recv from below)
102           cart_comm, MPI_STATUS_IGNORE);
103   // Calculation of matrix multiplication
104   c_ij += a_ij * b_ij;
105 }
106
107 // allocate memory for C matrix
108 float *C = (float *) calloc(sizeof(float), row * row);
109 // Gather the multiplication result from every processor
110 MPI_Gather(&c_ij, 1, MPI_FLOAT,
111                C, 1, MPI_FLOAT,
```

```
112             Master, cart_comm);
113
114   // Printing the result of Matrix multiplication stored in C array
115   if (rank == Master) {
116     int k = 0;
117     printf("\nA * B:\n");
118     for (int i = 0; i < row; i++) {
119       for (int j = 0; j < column; j++) {
120         printf("%f\t", C[k]);
121         k++;
122       }
123       printf("\n");
124     }
125   }
126
127   MPI_Finalize();
128   return 0;
129 }
130
131 // Assumes A and B will have same numberof rows and columns
132 //  TODO Should only assume that columns(A) == rows(B).
133 //      A_mxl . B_lxn  ==>  C_mxn
134 mat_size_t get_matrices(float **L, float **R) {
135   int row = 0;
136   int column = 0;
137   // finding the number of rows & columns in A matrix
138   FILE *fp;
139   fp = fopen("A.txt", "r");
140   int count = 0;
141   char ch;
142   float n;
143   // scan each line
144   while (fscanf(fp, "%f", &n) != -1) {
145     ch = fgetc(fp);
146     if (ch == '\n') {
147       row++;
148     }
149     count++;
150   }
151   column = count / row;
152
153   // Check to see to have enough processors for the elements
154   int cluster_size;
155   MPI_Comm_size(MPI_COMM_WORLD, &cluster_size);
156   if (count != cluster_size) {
157     printf("No of Processors must be equal to %d\nCode terminated\n", count);
158     MPI_Finalize();
159     fclose(fp);
160     exit(1);
161   }
162
163   // Jump back to beginning of file for matrix A
164   fseek(fp, 0, SEEK_SET);
165
166   // allocate memory for A and B
167   float *A = (float *)calloc(sizeof(float), row * column);
```

```
168   float *B = (float *)calloc(sizeof(float), row * column);
169
170   // Scanning and printing Matrix A
171   int k = 0;
172   printf("A matrix:\n");
173   for (int i = 0; i < row; i++) {
174     for (int j = 0; j < column; j++) {
175       fscanf(fp, "%f", &n);
176       A[k] = n;
177       printf("%f\t", A[k]);
178       k++;
179     }
180     printf("\n");
181   }
182   fclose(fp);
183
184   // Scanning and printing Matrix B
185   k = 0;
186   printf("\nB matrix:\n");
187   // read data for B matrix
188   fp = fopen("B.txt", "r");
189   for (int i = 0; i < row; i++) {
190     for (int j = 0; j < column; j++) {
191       fscanf(fp, "%f", &n);
192       B[k] = n;
193       printf("%f\t", B[k]);
194       k++;
195     }
196     printf("\n");
197   }
198   fclose(fp);
199
200   *L = A;
201   *R = B;
202   return (mat_size_t){row, column};
203 }
```